

RL Lab Assignment - 3 (Value Iteration and Policy Iteration)

Name : Arvind Pandit

Roll No. : 211022001

Value Iteration Algorithm

Initialize V arbitrarily, e.g., $V(s) = 0$, for all $s \in \mathcal{S}^+$

Repeat

$\Delta \leftarrow 0$

For each $s \in \mathcal{S}$:

$v \leftarrow V(s)$

$V(s) \leftarrow \max_a \sum_{s'} \mathcal{P}_{ss'}^a [\mathcal{R}_{ss'}^a + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

until $\Delta < \theta$ (a small positive number)

Output a deterministic policy, π , such that

$$\pi(s) = \arg \max_a \sum_{s'} \mathcal{P}_{ss'}^a [\mathcal{R}_{ss'}^a + \gamma V(s')]$$


Policy Iteration Algorithm

1. Initialization
 $V(s) \in \Re$ and $\pi(s) \in \mathcal{A}(s)$ arbitrarily for all $s \in \mathcal{S}$
2. Policy Evaluation
Repeat
 $\Delta \leftarrow 0$
 For each $s \in \mathcal{S}$:
 $v \leftarrow V(s)$
 $V(s) \leftarrow \sum_{s'} \mathcal{P}_{ss'}^{\pi(s)} [\mathcal{R}_{ss'}^{\pi(s)} + \gamma V(s')]$
 $\Delta \leftarrow \max(\Delta, |v - V(s)|)$
until $\Delta < \theta$ (a small positive number)
3. Policy Improvement
 policy-stable \leftarrow *true*
 For each $s \in \mathcal{S}$:
 $b \leftarrow \pi(s)$
 $\pi(s) \leftarrow \arg \max_a \sum_{s'} \mathcal{P}_{ss'}^a [\mathcal{R}_{ss'}^a + \gamma V(s')]$
 If $b \neq \pi(s)$, then *policy-stable* \leftarrow *false*
 If *policy-stable*, then stop; else go to 2

Reference: [Value iteration Policy Iteration](#)

Problem 1.

Here we have taken a very simple problem of Markov Decision Process.



There are five 2D tiles which represent the state of MDP.

State = {0,1,2,3,4,5}

The reward of each state is given by [-1, -1, 10, -1, -1] and the goal is to reach terminal state 3rd which have reward of 10.

Reward = {-1, -1, 10, -1, -1}

The allowed action are left and right

Action = {0,1} where 0=left and 1=right

Here the agent have to collect the maximum reward and reach to the terminal state 3rd.

Probability Transition matrix for action a1-left

```
[[0.9 0.1 0.  0.  0.  ]
 [0.9 0.  0.1 0.  0.  ]
 [0.  0.  0.  0.  0.  ]
 [0.  0.  0.9 0.  0.1]
 [0.  0.  0.  0.9 0.1]]
```

Probability Transition matrix for action a2-right

```
[[0.1 0.9 0.  0.  0.  ]
 [0.1 0.  0.9 0.  0.  ]
 [0.  0.  0.  0.  0.  ]
 [0.  0.  0.1 0.  0.9]
 [0.  0.  0.  0.1 0.9]]
```

1. Value Iteration

No. of iterations in Value Iterations = 10

For Value iteration

Optimal Value function $V^* = [5.67496513 \quad 7.61064654 \quad 10. \quad 7.61064654 \quad 5.67496513]$

Optimal policy $\mu^* = [1, 1, 0, 0, 0]$

2. Policy Iterations


No. of iterations in Policy Iterations = 2

For Policy iteration

Optimal Value function $V^* = [5.67450141 \quad 7.61052229 \quad 10. \quad 7.61079919 \quad 5.67554653]$

Optimal policy $\mu^* = [1. \quad 1. \quad 0. \quad 0. \quad 0.]$

The Optimal Policy is $\mu = [1 \ 1 \ 0 \ 0 \ 0]$ s.t [right right left left left]



3. Verification by choosing Random Policy

Average collected reward with optimal policy μ [1. 1. 0. 0. 0.] is 7.314273885011818

Average collected reward with random policy μ_0 [0 1 1 0 1] is 3.9819028817869593

Average collected reward with random policy μ_1 [1 1 0 0 1] is 5.645207726534246


Average collected reward with random policy μ_2 [1 1 1 1 0] is 3.104047648468936

Problem 2.

We have taken a grid problem of Markov Decision Process.

Here the robot starts from the state 0 and there are walls in state (2,3,8,9) and the final destination of the robot is state 12.

The movement of the robot have the randomness such that it will go to the commanded direction with probability 0.7 and other 3 direction with each having probability 0.1.



There are 16 grids which represent the state of MDP.

State = {0,1,2,3,4,5,6,7,8,9,10,11,12,13,14,15}

The reward of each state is given by {-1,-1,0,0,-1,-1,-1,-1,0,0,-1,-1,10,-1,-1,-1} and the goal is to reach terminal state 12th which have reward of 10.

Each step will cost 1.

Reward = {-1,-1,0,0,-1,-1,-1,-1,0,0,-1,-1,10,-1,-1,-1}

The allowed action are up,down,left and right

Action = {0,1,2,3} where 0=up,1=down,2=left,3=right

Here the agent have to collect the maximum reward and reach to the terminal state 12 avoiding the collision with walls.

Probability Transition matrix for action a1-up

```
[
[0.2 0.7 0. 0. 0.1 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. ]
[0.1 0.8 0. 0. 0. 0.1 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. ]
[0. 0. 1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. ]
[0. 0. 0. 1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. ]
[0.1 0. 0. 0. 0. 0.1 0.7 0. 0. 0.1 0. 0. 0. 0. 0. 0. ]
[0. 0.1 0. 0. 0.1 0. 0.7 0. 0. 0.1 0. 0. 0. 0. 0. 0. ]
[0. 0. 0.1 0. 0. 0.1 0. 0.7 0. 0. 0.1 0. 0. 0. 0. 0. ]
[0. 0. 0. 0.1 0. 0. 0.1 0.7 0. 0. 0. 0.1 0. 0. 0. 0. ]
[0. 0. 0. 0. 0. 0. 0. 0. 1. 0. 0. 0. 0. 0. 0. 0. 0. ]
[0. 0. 0. 0. 0. 0. 0. 0. 0. 1. 0. 0. 0. 0. 0. 0. 0. ]
[0. 0. 0. 0. 0. 0. 0.1 0. 0. 0.1 0. 0.7 0. 0. 0.1 0. 0. ]
[0. 0. 0. 0. 0. 0. 0. 0.1 0. 0. 0.1 0.7 0. 0. 0. 0.1 ]
[0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 1. 0. 0. 0. 0. ]
]
```

```
[0.  0.  0.  0.  0.  0.  0.  0.  0.  0.1 0.  0.  0.1 0.1 0.7 0. ]
[0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.1 0.  0.  0.1 0.1 0.7]
[0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.1 0.  0.  0.1 0.8]]
```

Probability Transition matrix for action a2-down

```
[[0.8 0.1 0.  0.  0.1 0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0. ]
 [0.7 0.1 0.1 0.  0.  0.1 0.  0.  0.  0.  0.  0.  0.  0.  0.  0. ]
 [0.  0.7 0.1 0.1 0.  0.  0.1 0.  0.  0.  0.  0.  0.  0.  0.  0. ]
 [0.  0.  0.  0.9 0.  0.  0.  0.1 0.  0.  0.  0.  0.  0.  0.  0. ]
 [0.1 0.  0.  0.  0.7 0.1 0.  0.  0.1 0.  0.  0.  0.  0.  0.  0. ]
 [0.  0.1 0.  0.  0.7 0.  0.1 0.  0.  0.1 0.  0.  0.  0.  0.  0. ]
 [0.  0.  0.1 0.  0.  0.7 0.  0.1 0.  0.  0.1 0.  0.  0.  0.  0. ]
 [0.  0.  0.  0.1 0.  0.  0.7 0.1 0.  0.  0.  0.1 0.  0.  0.  0. ]
 [0.  0.  0.  0.  0.1 0.  0.  0.  0.7 0.1 0.  0.  0.1 0.  0.  0. ]
 [0.  0.  0.  0.  0.  0.1 0.  0.  0.  0.7 0.1 0.  0.  0.1 0.  0. ]
 [0.  0.  0.  0.  0.  0.  0.1 0.  0.  0.  0.7 0.1 0.  0.  0.1 0. ]
 [0.  0.  0.  0.  0.  0.  0.  0.1 0.  0.  0.  0.7 0.1 0.  0.  0.1]
 [0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  1.  0.  0.  0. ]
 [0.  0.  0.  0.  0.  0.  0.  0.  0.  0.1 0.  0.  0.7 0.1 0.1 0. ]
 [0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.1 0.  0.  0.7 0.1 0.1]
 [0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.1 0.  0.  0.7 0.2]]
```

Probability Transition matrix for action a3-left

```
[[0.8 0.1 0.  0.  0.1 0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0. ]
 [0.1 0.7 0.1 0.  0.  0.1 0.  0.  0.  0.  0.  0.  0.  0.  0.  0. ]
 [0.  0.  1.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0. ]
 [0.  0.  0.  1.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0. ]
 [0.7 0.  0.  0.  0.1 0.1 0.  0.  0.1 0.  0.  0.  0.  0.  0.  0. ]
 [0.  0.7 0.  0.  0.1 0.  0.1 0.  0.  0.1 0.  0.  0.  0.  0.  0. ]
 [0.  0.  0.  0.  0.  0.1 0.7 0.1 0.  0.  0.1 0.  0.  0.  0.  0. ]
 [0.  0.  0.  0.  0.  0.  0.1 0.8 0.  0.  0.  0.1 0.  0.  0.  0. ]
 [0.  0.  0.  0.  0.  0.  0.  0.  1.  0.  0.  0.  0.  0.  0.  0. ]
 [0.  0.  0.  0.  0.  0.  0.  0.  0.  1.  0.  0.  0.  0.  0.  0. ]
 [0.  0.  0.  0.  0.  0.  0.7 0.  0.  0.1 0.  0.1 0.  0.  0.1 0. ]
 [0.  0.  0.  0.  0.  0.  0.  0.7 0.  0.  0.1 0.1 0.  0.  0.  0.1]
 [0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  1.  0.  0.  0. ]
 [0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.1 0.8 0.1 0. ]
 [0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.7 0.  0.  0.1 0.1 0.1]
 [0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.7 0.  0.  0.1 0.2]]
```

Probability Transition matrix for action a4-right

```
[[0.2 0.1 0.  0.  0.7 0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0. ]
 [0.1 0.1 0.1 0.  0.  0.7 0.  0.  0.  0.  0.  0.  0.  0.  0.  0. ]
```

```
[0.  0.  1.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0. ]
[0.  0.  0.  1.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0. ]
[0.1 0.  0.  0.  0.8 0.1 0.  0.  0.  0.  0.  0.  0.  0.  0.  0. ]
[0.  0.1 0.  0.  0.1 0.7 0.1 0.  0.  0.  0.  0.  0.  0.  0.  0. ]
[0.  0.  0.1 0.  0.  0.1 0.  0.1 0.  0.  0.7 0.  0.  0.  0.  0. ]
[0.  0.  0.  0.1 0.  0.  0.1 0.1 0.  0.  0.  0.7 0.  0.  0.  0. ]
[0.  0.  0.  0.  0.  0.  0.  0.  1.  0.  0.  0.  0.  0.  0.  0. ]
[0.  0.  0.  0.  0.  0.  0.  0.  0.  1.  0.  0.  0.  0.  0.  0. ]
[0.  0.  0.  0.  0.  0.  0.1 0.  0.  0.1 0.  0.1 0.  0.  0.7 0. ]
[0.  0.  0.  0.  0.  0.  0.  0.1 0.  0.  0.1 0.1 0.  0.  0.  0.7]
[0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  1.  0.  0.  0. ]
[0.  0.  0.  0.  0.  0.  0.  0.  0.  0.1 0.  0.  0.1 0.7 0.1 0. ]
[0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.1 0.  0.  0.1 0.7 0.1]
[0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.1 0.  0.  0.1 0.8]]
```

1. Value Iteration

No. of iterations in Value Iterations = 76

For Value iteration

Optimal Value function $V^* = [21.89782805 \ 24.80772769 \ 22.53287879$

$16.14270632 \ 26.55055866 \ 31.07872535$

$38.0726069 \ 34.11590598 \ 40.15483901 \ 38.60073338 \ 49.49279352$

43.12606991

$99.96670104 \ 78.15737539 \ 63.03837479 \ 51.94169849]$

Optimal policy $\mu^* = [3, 3, 1, 1, 0, 0, 3, 3, 1, 1, 3, 3, 0, 1, 1, 1]$

Optimal policy['right', 'right', 'nan', 'nan', 'up', 'up', 'right',
'right', 'nan', 'nan', 'right', 'right', 'nan', 'down', 'down',
'down']

2. Policy Iteration

$[3. \ 3. \ 0. \ 0. \ 0. \ 0. \ 1. \ 1. \ 1. \ 0. \ 2. \ 2. \ 0. \ 1. \ 1. \ 1.]$

$[3. \ 1. \ 0. \ 0. \ 1. \ 1. \ 3. \ 3. \ 1. \ 1. \ 3. \ 3. \ 0. \ 1. \ 1. \ 1.]$

$[3. \ 3. \ 1. \ 1. \ 0. \ 0. \ 3. \ 3. \ 1. \ 1. \ 3. \ 3. \ 0. \ 1. \ 1. \ 1.]$

$[3. \ 3. \ 1. \ 1. \ 0. \ 0. \ 3. \ 3. \ 1. \ 1. \ 3. \ 3. \ 0. \ 1. \ 1. \ 1.]$

No. of iterations in Policy Iterations = 4

For Value iteration

Optimal Value function $V^* = [\ 21.92829068 \ 24.83823287 \ 22.5612402$

$16.15729915 \ 26.58141653$



$31.10992607 \ 38.10416901 \ 34.14586757 \ 40.18609368 \ 38.63246463$

$49.52532705 \ 43.15841139 \ 100. \ 78.19043061 \ 63.07129906$

$51.9744841]$

Optimal policy $\mu^* = [3. \ 3. \ 1. \ 1. \ 0. \ 0. \ 3. \ 3. \ 1. \ 1. \ 3. \ 3. \ 0. \ 1. \ 1. \ 1.]$

Optimal policy['right', 'right', 'nan', 'nan', 'up', 'up', 'right',
'right', 'nan', 'nan', 'right', 'right', 'nan', 'down', 'down',
'down']

3	7 →	11 →	15 ↓
2	6 →	10 →	14 ↓
1 →	5 ↑	9	13 ↓
0 → 	4 ↑	8	12 

3. Verification by choosing Random Policy

Verification by choosing Random Policy

Average collected reward with optimal policy μ [3. 3. 1. 1. 0. 0. 3. 3. 1. 1. 3. 3. 0. 1. 1. 1.] is 42.51030953722727

Average collected reward with random policy μ [2 3 2 1 0 0 2 2 0 3 0 1 0 2 1 3] is 7.855080063370282

Average collected reward with random policy μ [2 3 0 1 0 2 2 3 0 1 0 0 1 2 1 3] is 8.715658344714956

Average collected reward with random policy μ [3 0 2 1 3 3 0 3 2 0 3 3 3 2 2 3] is 3.951926359712211

Average collected reward with random policy μ [3 1 1 3 0 2 2 0 0 1 0 0 2 2 3 3] is 4.679206897793618

Average collected reward with random policy μ [3 1 3 0 3 2 2 1 3 1 1 1 2 2 3 2] is 4.390354843189319

4. Observations

Problem-1

Optimal Value function $V^* = [5.67 \ 7.61 \ 10. \ 7.61 \ 5.67]$

Optimal Policy $\mu^* = ['right', 'right', 'nan', 'left', 'left']$

Value Iteration

No. of iterations in Value Iterations = 10

Policy Iteration

No. of iterations for $V_{\mu_0} = 49$

No. of iterations for $V_{\mu_1} = 9$

No. of iterations in Policy Iterations = 2

Problem-2

Optimal Value function $V^* = [21.89 \ 24.80 \ 22.53 \ 16.14 \ 26.55 \ 31.07 \ 38.07 \ 34.11 \ 40.15 \ 38.60 \ 49.49 \ 43.12 \ 99.96 \ 78.15 \ 63.03 \ 51.94]$

Optimal policy $\mu^* = ['right', 'right', 'nan', 'nan', 'up', 'up', 'right', 'right', 'nan', 'nan', 'right', 'right', 'nan', 'down', 'down', 'down']$

Value Iteration

No. of iterations in Value Iterations = 76

Policy Iteration

No. of iterations for $V_{\mu_0} = 67$

No. of iterations for $V_{\mu_1} = 28$

No. of iterations for $V_{\mu_2} = 45$

No. of iterations for $V_{\mu_3} = 36$

No. of iterations in Policy Iterations = 4